

# Managing Data in Docker

**Prepared by:**

Ms. Avita Katal

Assistant Professor (SG)

School of Computer Science

UPES, Dehradun

# Objectives

- What do you understand by managing data in Docker
- Learn about different ways of managing data and their use cases
  - Volumes
  - Bind mounts
  - tmpfs mount

# Manage data in Docker

By default all files created inside a container are stored on a **writable container layer**. This means that:

- The data **doesn't persist when that container no longer exists**, and it can be difficult to get the data out of the container if another process needs it.
- A container's **writable layer is tightly coupled to the host machine** where the container is running. You can't easily move the data somewhere else.
- Writing into a **container's writable layer requires a storage driver to manage the file system**. The storage driver provides a **union file system**, using the Linux kernel. This extra abstraction reduces performance as compared to using **data volumes**, which **write directly to the host file system**.

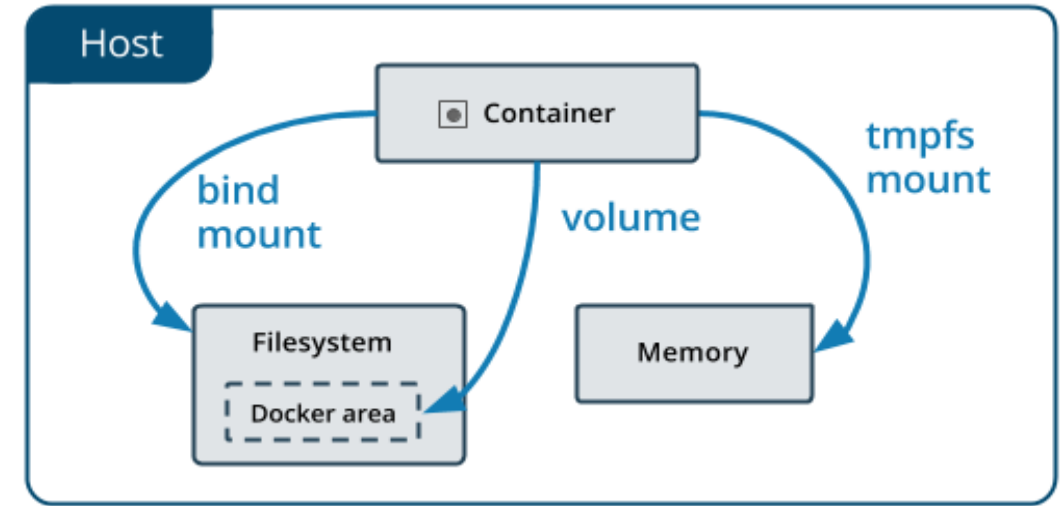
Docker has two options for **containers to store files on the host machine**, so that the files are persisted even after the container stops:

*volumes, and bind mounts.*

Docker also supports containers **storing files in-memory on the host machine**. Such files are **not persisted**. If you're running **Docker on Linux**, *tmpfs mount* is used to store files in the host's system memory. If you're running **Docker on Windows**, *named pipe* is used to store files in the host's system memory.

# Manage data in Docker

- No matter which type of mount you choose to use, the data looks the same from within the container. It is exposed **as either a directory or an individual file in the container's file system.**
- An easy way to visualize the difference among **volumes**, **bind mounts**, and **tmpfs mounts** is to think about where the data lives on the **Docker host**.



- **Volumes** are stored in a part of the **host file system** which is *managed by Docker* (**`/var/lib/docker/volumes/`** on **Linux**). **Non-Docker processes should not modify this part of the file system.** Volumes are the best way to persist data in Docker.
- **Bind mounts** may be stored *anywhere* on the **host system**. They may even be important system files or directories. **Non-Docker processes on the Docker host or a Docker container can modify them at any time.**
- **tmpfs mounts** are stored in the **host system's memory only**, and are **never written to the host system's filesystem.**

# Docker Volumes

- Created and managed by Docker. You can create a volume explicitly using the ***docker volume create command***, or Docker can **create a volume during container or service creation**.
- When you create a volume, it is **stored within a directory on the Docker host**. When you mount the volume into a container, this directory is what is mounted into the container. This is similar to the way that bind mounts work, except that **volumes are managed by Docker and are isolated from the core functionality of the host machine**.
- A given volume **can be mounted into multiple containers** simultaneously. When no running container is using a volume, the **volume is still available to Docker and is not removed automatically**.
- You can remove unused volumes using ***docker volume prune***.
- When you mount a volume, it may be **named or anonymous**.
  - Anonymous volumes **are not given an explicit name** when they are first mounted into a container, so Docker gives **them a random name** that is guaranteed to be unique within a given Docker host. Besides the name, named and anonymous volumes behave in the same ways.
- Volumes also support the use of **volume drivers**, which allow you to store your data on **remote hosts or cloud providers**, among other possibilities.

# Good use cases for volumes

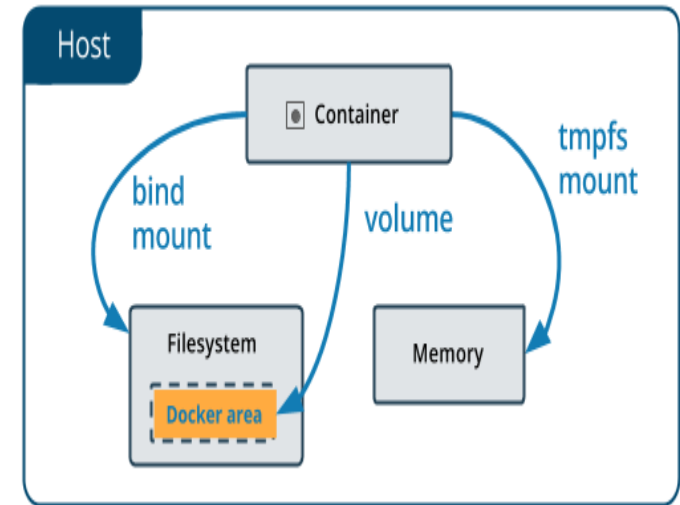
- **Sharing data among multiple running containers.** Multiple containers can mount the same volume simultaneously, either read-write or read-only. Volumes are only removed when you explicitly remove them.
- When the Docker host is not guaranteed to have a given directory or file structure. Volumes help you **decouple the configuration of the Docker host from the container runtime.**
- When you want to store your **container's data on a remote host or a cloud provider**, rather than locally.
- When you need to **back up, restore, or migrate data from one Docker host to another**, volumes are a better choice. You can stop containers using the volume, then back up the volume's directory (such as `/var/lib/docker/volumes/<volume-name>`).
- When **your application requires high-performance I/O on Docker Desktop**. Volumes are stored in the **Linux VM rather than the host**, which means that the reads and writes have much lower latency and higher throughput.
- When **your application requires fully native file system** behavior on Docker Desktop.

# Docker Volumes

Volumes are the preferred mechanism for persisting data generated by and used by Docker containers. **While bind mounts are dependent on the directory structure and OS of the host machine, volumes are completely managed by Docker.** Volumes have several advantages over bind mounts:

- Volumes are *easier to back up or migrate* than bind mounts.
- You can manage *volumes using Docker CLI commands or the Docker API.*
- Volumes work on *both Linux and Windows containers.*
- Volumes can be *more safely shared among multiple containers.*
- Volume drivers let you *store volumes on remote hosts or cloud providers, to encrypt the contents of volumes,* or to add other functionality.
- New volumes can have *their content pre-populated by a container.*
- Volumes on *Docker Desktop have much higher performance than bind mounts* from Mac and Windows hosts.

In addition, *volumes are often a better choice than persisting data in a container's writable layer, because a volume doesn't increase the size of the containers using it, and the **volume's contents exist outside the lifecycle of a given container.***



## Create and manage volumes

Unlike a bind mount, **you can create and manage volumes outside the scope of any container.**

### Create a volume:

```
docker volume create my-vol
```

### List volumes:

```
docker volume ls
```

### Inspect a volume:

```
docker volume inspect my-vol
```

### Remove a volume:

```
docker volume rm my-vol
```



# Start a container with a volume

If you **start a container with a volume that doesn't yet exist, Docker creates the volume for you**. The following example mounts the volume myvol2 into /app/ in the container.

The **-v** and **--mount** examples below produce the same result. You can't run them both unless you remove the devtest container and the myvol2 volume after running the first one.

```
docker run -d \  
  --name devtest \  
  --mount source=myvol2,target=/app \  
  nginx:latest
```

```
docker run -d \  
  --name devtest \  
  -v myvol2:/app \  
  nginx:latest
```

Use **docker inspect devtest** to verify that Docker created the volume and it mounted correctly.

**Stop the container and remove the volume. Note volume removal is a separate step.**

```
docker container stop devtest  
docker container rm devtest  
docker volume rm myvol2
```

# Bind Mounts

- Available since the **early days of Docker**. Bind mounts have **limited functionality** compared to volumes.
- When you use a bind mount, **a file or directory on the *host machine* is mounted into a container.**
- The file or directory is **referenced by its full path on the host machine.**
- The file or directory **does not need to exist on the Docker host already.**
- It is created **on demand if it does not yet exist**. Bind mounts are **very performant, but they rely on the host machine's file system** having a specific directory structure available.
- If you are developing new Docker applications, consider using **named volumes instead**. **You can't use Docker CLI commands to directly manage bind mounts.**

## *Bind mounts allow access to sensitive files*

One side effect of using bind mounts, for better or for worse, is that you can change the **host** file system via processes running in a **container**, including creating, modifying, or deleting important system files or directories. This is a powerful ability which can have ***security implications***, including impacting **non-Docker processes on the host system**.

# Good use cases for bind mount

- **Sharing configuration files from the host machine to containers.** This is how Docker provides **DNS resolution to containers** by default, by mounting `/etc/resolv.conf` from the host machine into each container.
- **Sharing source code or build artifacts between a development environment on the Docker host and a container.** For instance, you may mount a Maven `target/` directory into a container, and each time you build the Maven project on the Docker host, the container gets access to the rebuilt artifacts. **If you use Docker for development this way, your production Dockerfile would copy the production-ready artifacts directly into the image, rather than relying on a bind mount.**
- When the **file or directory structure of the Docker host is guaranteed to be consistent with the bind mounts the containers require.**

# Start a container with bind mount

- Consider a case where you have a directory source and that when you build the source code, the artifacts are saved into another directory, **source/target/**.
- You want the artifacts to be available to the container at **/app/**, and you want the container to get access to a new build each time you build the source on your development host.
- Use the following command to bind-mount the **target/ directory** into your container at **/app/**. Run **the command from within the source directory**.
- The **\$(pwd)** sub-command expands to the current working directory on Linux or macOS hosts. If you're on Windows, see also Path conversions on Windows.

The --mount and -v examples below produce the same result. You can't run them both unless you remove the devtest container after running the first one.

```
$ docker run -d \  
-it \  
--name devtest \  
-v "$(pwd)"/target:/app \  
nginx:latest
```

```
$ docker run -d \  
-it \  
--name devtest \  
--mount type=bind, source="$(pwd)"/target, target=/app \  
nginx:latest
```

*docker inspect devtest*

to verify that the bind mount was created correctly. Look for the Mounts section.

**Stop the container:**

\$ docker container stop devtest

\$ docker container rm devtest

# Tmps Mounts and Named Pipes

## tmpfs mounts

- A **tmpfs** mount is **not persisted on disk**, either on the Docker host or within a container.
- It can be used by a **container during the lifetime of the container**, to store **non-persistent state or sensitive information**.
- For instance, internally, swarm services use **tmpfs** mounts **to mount secrets into a service's containers**.

## named pipes

- An **npipe** mount can be used for **communication between the Docker host and a container**.
- Common use case is to **run a third-party tool inside of a container and connect to the Docker Engine API using a named pipe**.



# Good use cases for tmpfs mounts

- tmpfs mounts are best used for cases when you do **not want the data to persist either on the host machine or within the container.**
- This may be for **security reasons or to protect the performance of the container when your application needs to write a large volume of non-persistent state data.**

## Limitations for tmpfs mounts

- Unlike volumes and bind mounts, you **can't share tmpfs mounts between containers.**
- This functionality is only available **if you're running Docker on Linux.**
- Setting **permissions on tmpfs may cause them to reset after container restart.**

# Use tmpfs mount in a container

- To use a tmpfs mount in a container, use the **--tmpfs flag**, or use the **--mount flag** with **type=tmpfs** and destination options.
- There is no source for **tmpfs mounts**. The following example creates a **tmpfs mount at /app** in a Nginx container.
- The first example uses the **--mount flag** and the second uses the **--tmpfs flag**.

```
docker run -d \  
-it \  
--name tmptest \  
--mount type=tmpfs,destination=/app \  
nginx:latest
```

```
docker run -d \  
-it \  
--name tmptest \  
--tmpfs /app \  
nginx:latest
```

**Verify that the mount is a tmpfs mount by looking in the Mounts section of the docker inspect output:**

**Stop and remove the container:**

```
$ docker stop tmptest
```

```
$ docker rm tmptest
```



**Thank You**